

Before you turn this problem in, make sure everything runs as expected. First, **restart the kernel** (in the menubar, select Kernel→Restart) and then **run all cells** (in the menubar, select Cell→Run All).

Make sure you fill in any place that says `YOUR CODE HERE` or "YOUR ANSWER HERE", as well as your name and collaborators below:

```
In [ ]: NAME = ""
        COLLABORATORS = ""
```

```
In [ ]: try:
        import torch as t
        import torch.nn as tnn
    except ImportError:
        print("Colab users: pytorch comes preinstalled. Select Change Ru")
        print("Local users: Please install pytorch for your hardware using instr")
        print("ACG users: Please follow instructions here: https://vikasdhiman.i")

        raise
```

```
In [ ]: def wget(url, filename):
        """
        Download files using requests package.
        Better than wget command line because this is cross platform.
        """
        try:
            import requests
        except ImportError:
            import subprocess
            subprocess.call("pip install --user requests".split())
            import requests
        r = requests.get(url)
        with open(filename, 'wb') as fd:
            for chunk in r.iter_content():
                fd.write(chunk)
```

```
In [ ]: # Get training features from MNIST dataset.
        wget("https://vikasdhiman.info/ECE490-Neural-Networks/notebooks/05-mlp/zero_
            "zero_one_train_features.npz")
```

```
In [ ]: def draw_features(ax, zero_features, one_features):
        zf = ax.scatter(zero_features[:, 0], zero_features[:, 1], marker='.', la
        of = ax.scatter(one_features[:, 0], one_features[:, 1], marker='+', labe
        ax.legend()
        ax.set_xlabel('Feature 1: count of pixels')
        ax.set_ylabel('Feature 2: Variance along x-axis')
        return [zf, of] # return list of artists
```

```
In [ ]: import numpy as np
        import matplotlib.pyplot as plt
```

```

zero_one_train_features = np.load('zero_one_train_features.npz')
FEATURE_MEAN = zero_one_train_features['mean']
FEATURE_STD = zero_one_train_features['std']
features = zero_one_train_features['normed_features']
labels = zero_one_train_features['labels']

fig, ax = plt.subplots()
draw_features(ax, features[labels > 0, :], features[labels < 0, :])

```

```

In [ ]: if t.cuda.is_available():
        DEVICE="cuda"
    elif t.mps.is_available():
        DEVICE="mps"
    else:
        DEVICE="cpu"

    DTYPE = t.get_default_dtype()

    def loss(predicted_labels, true_labels):
        # Make sure predicted_labels and true_labels have same shape
        y = true_labels[..., None]
        yhat = predicted_labels
        assert y.shape == yhat.shape
        return t.maximum(- y * yhat, t.Tensor([0.]).to(device=DEVICE)).sum() / y

    # TODO:
    # Define model = ?
    model = tnn.Sequential(
        tnn.Linear(2, 5),
        tnn.ReLU(),
        tnn.Linear(5, 1))

    def train_by_gradient_descent(model, loss, train_features, train_labels, lr=
        predicted_labels = model(train_features)
        #print(predicted_labels)

        loss_t = loss(predicted_labels, train_labels)
        loss_t.backward()
        loss_t_minus_1 = 2*loss_t # Fake value to make the while test pass once
        niter = 0
        while t.abs(loss_t - loss_t_minus_1) / loss_t > 0.01: # Stopping criteria
            with t.no_grad(): # parameter update needs no gradients
                for param in model.parameters():
                    assert param.grad is not None
                    param.add_(- lr * param.grad) # Gradient descent

            model.zero_grad()
            # Recompute the gradients
            predicted_labels = model(train_features)
            loss_t_minus_1 = loss_t
            loss_t = loss(predicted_labels, train_labels)
            loss_t.backward() # Compute gradients for next iteration

            # If loss increased, decrease lr. Works for gradient descent, not for
            if loss_t > loss_t_minus_1:

```

```

        lr = lr / 2

    ### DEBUGging information
    iswrong = (train_labels * predicted_labels.ravel()) < 0
    misclassified = (iswrong).sum() / iswrong.shape[0]
    print(f"loss: {loss_t:04.04f}, delta loss: {loss_t - loss_t_minus_1:04.04f}, "
          f"train misclassified: {misclassified:04.04f}")
    if niter % 20 == 0: # plot every 20th iteration
        train_features_cpu = train_features.cpu()
        predicted_labels_cpu = predicted_labels.cpu()
        fig, ax = plt.subplots(1,1)
        draw_features(ax,
                      train_features_cpu[predicted_labels_cpu.ravel() > 0, :],
                      train_features_cpu[predicted_labels_cpu.ravel() < 0, :])

    niter += 1
    return model

trained_model = train_by_gradient_descent(model.to(device=DEVICE),
                                           loss,
                                           t.from_numpy(features).to(device=DEVICE),
                                           t.from_numpy(labels).to(device=DEVICE))

fig, axes = plt.subplots(1,2)
draw_features(axes[0], features[labels > 0, :], features[labels < 0, :])
axes[0].set_title('Train labels')

predicted_labels = trained_model(t.from_numpy(features).to(device=DEVICE, dtype=torch.float))
predicted_labels_cpu = predicted_labels.cpu()
draw_features(axes[1], features[predicted_labels_cpu.ravel() > 0, :],
              features[predicted_labels_cpu.ravel() < 0, :])
axes[1].set_title('Predicted labels');

```

In []: *## Doing it the Pytorch way without using our custom feature extraction*

```

import torch
import torch.nn
import torch.optim
import torchvision
from torchvision.transforms import ToTensor
from torch.utils.data import DataLoader

torch.manual_seed(17)

# Getting the dataset, the Pytorch way
all_training_data = torchvision.datasets.MNIST(
    root="data",
    train=True,
    download=True,
    transform=ToTensor()
)

test_data = torchvision.datasets.MNIST(

```

```

    root="data",
    train=False,
    download=True,
    transform=ToTensor()
)

```

```
In [ ]: training_data, validation_data = torch.utils.data.random_split(all_training_
```

```
In [ ]: # Hyper parameters
learning_rate = 1e-3 # controls how fast the
batch_size = 64
epochs = 5
momentum = 0.9

training_dataloader = DataLoader(training_data, shuffle=True, batch_size=batch_size)
validation_dataloader = DataLoader(validation_data, batch_size=batch_size)
test_dataloader = DataLoader(test_data, batch_size=batch_size)

loss = torch.nn.CrossEntropyLoss()
# TODO:
# Define model = ?

model = tnn.Sequential(
    torch.nn.Flatten(),
    tnn.Linear(28*28, 10),
    tnn.ReLU(),
    tnn.Linear(10, 10))

# Define optimizer
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate, momentum=0.9)

def loss_and_accuracy(model, loss, validation_dataloader, device=DEVICE):
    # Validation loop
    validation_size = len(validation_dataloader.dataset)
    num_batches = len(validation_dataloader)
    test_loss, correct = 0, 0

    with torch.no_grad():
        for X, y in validation_dataloader:
            X = X.to(device)
            y = y.to(device)
            pred = model(X)
            test_loss += loss(pred, y).item()
            correct += (pred.argmax(dim=-1) == y).type(DTYPE).sum().item()

    test_loss /= num_batches
    correct /= validation_size
    return test_loss, correct

def train(model, loss, training_dataloader, validation_dataloader, device=DEVICE):
    model.to(device)
    for t in range(epochs):
        # Train loop
        training_size = len(training_dataloader.dataset)

```

```

    for batch, (X, y) in enumerate(training_dataloader):
        X = X.to(device)
        y = y.to(device)
        # Compute prediction and loss
        pred = model(X)
        loss_t = loss(pred, y)

        # Backpropagation
        optimizer.zero_grad()
        loss_t.backward()
        optimizer.step()

        if batch % 100 == 0:
            loss_t, current = loss_t.item(), (batch + 1) * len(X)
            print(f"loss: {loss_t:>7f} [{current:>5d}/{training_size:>5d}]")
            valid_loss, correct = loss_and_accuracy(model, loss, validation_data_loader)
            print(f"Validation Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss: {valid_loss:>0.1f}")
    return model

trained_model = train(model, loss, training_dataloader, validation_data_loader)

test_loss, correct = loss_and_accuracy(model, loss, test_dataloader)
print(f"Test Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss: {test_loss:>0.1f}")

```

```

In [ ]: X, _ = next(iter(test_dataloader))
        X.shape

```

```

In [ ]: import matplotlib.pyplot as plt
        plt.imshow(X[0, 0])

```

```

In [ ]: print("The predicted image label is ", model(X.to(DEVICE)).argmax(dim=-1)[0])

```