



Supervised Learning + NN

- ① Dataset $\{(x_1, y_1) \dots\}$
- ② Model $\begin{cases} \text{Linear model} \\ \text{MLP model} \end{cases}$
- ③ Loss $\begin{cases} \text{MSE} \\ \text{Perceptron loss} \\ \text{cross entropy loss} \end{cases}$
- ④ Optimization methods $\begin{cases} \text{GD} \\ \text{SGD} \end{cases}$


MNIST




CIFAR10



Flatten
→




Input



$+b_1 \rightarrow$



\rightarrow




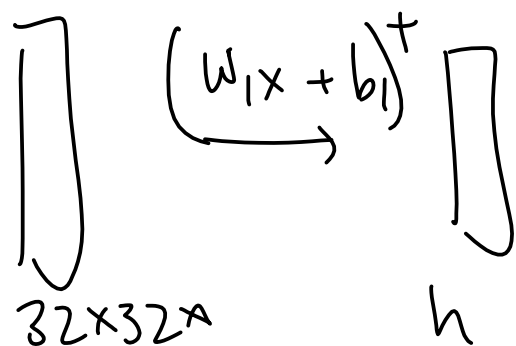
Output

g

(13) Hidden Layer

Size(HL) $\rightarrow \infty$

Num
classes
 $= 10$



Cross entropy loss

$$L(Y, \hat{y}) = \sum_{y=0}^9 -P(Y=y) \log \left(\frac{\exp(\hat{y}[y])}{\sum_{i=0}^9 \exp(\hat{y}[i])} \right)$$

$$P(\hat{Y}=y)$$

Y is a RV $\Omega_Y = \{0, 1, \dots, 9\}$ $P(\hat{Y}=y)$

$(-\infty, \infty)$

Output
of model

10
20
30
40
50
10

\hat{y}

$$\frac{\exp(\cdot)}{\sum \exp(\cdot)}$$

softmax


$\exp(10)$	0
$\exp(20)$	1
$\exp(30)$	2
\vdots	\vdots
$\exp(10)$	9

$(0, 1)$

hardmax $\rightarrow 40$

$$\frac{10^{\hat{y}}}{\sum 10^{\hat{y}}}$$


10^{10}
10^{20}
10^{30}
10^{40}
10^{50}
10^{10}



$$n = 1024$$

$$\underbrace{10^9 \times 4 \text{ bytes}} = 4 \text{ GB}$$

Images




Position of the tree in image classification problem does not matter

This property is called translation invariance

Convolution operation

that is translation invariant



4

0	1	0	0
1	1	1	0
0	1	0	0
0	0	0	0

4

Image

=

Convolution Kernel

3


0	1	0
0	1	0
0	1	0

3

3

3	1
2	1

$$\Sigma \textcircled{1} \times \textcircled{1} + \textcircled{2} \times \textcircled{2} + \textcircled{3} \times \textcircled{3} + \dots$$




out-channels \rightarrow create more kernels
 for 10 = out-channels
 create 10 convolutional kernels


in-channels

kernel size


Size is shape of
 conv kernel along
 image dimensions



stride = step size = 1




stride = step size = 2



padding

$$\begin{pmatrix} 4 \times 4 \\ \text{image} \end{pmatrix} * \begin{pmatrix} 3 \times 3 \\ \text{kernel} \end{pmatrix} = \begin{pmatrix} 2 \times 2 \\ \text{output} \end{pmatrix}$$




$$\begin{pmatrix} 4 \times 4 \\ \text{image} \end{pmatrix} \xrightarrow{\text{padding} \Rightarrow} \begin{pmatrix} 6 \times 6 \\ \text{image} \end{pmatrix} * \begin{pmatrix} 3 \times 3 \\ \text{kernel} \end{pmatrix} = \begin{pmatrix} 4 \times 4 \\ \text{output} \end{pmatrix}$$

$$\begin{pmatrix} 10 \times 10 \\ \text{image} \end{pmatrix} * \begin{pmatrix} 3 \times 3 \\ \text{kernel} \end{pmatrix} = \left(10 - \left\lceil \frac{3}{2} \right\rceil \times 2, 10 - \left\lfloor \frac{3}{2} \right\rfloor \times 2 \right)$$

$$= \begin{pmatrix} 8 \times 8 \\ \text{output size} \end{pmatrix}$$

A diagram of a 10x10 grid representing an image. The top-left 3x3 area is highlighted with a grid, representing the kernel operation.

Output size of convolution kernel




$$O_w = w - k_w + 1$$

$$O_h = h - k_h + 1$$

$$O_w = \text{floor} \left\lfloor \frac{w - k_w + 1}{S_w} \right\rfloor$$

$$O_h = \text{floor} \left\lfloor \frac{h - k_h + 1}{S_h} \right\rfloor$$

	i				
	1	1	1	1	1
	1	15	17	15	1
j	1	17	20	17	1
	1	15	17	15	1
	1	1	1	1	1





$$k(i,j) = \exp \left(- \frac{1}{2} \frac{(i - \mu_i)^2 + (j - \mu_j)^2}{\sigma_i^2 \sigma_j^2} \right) \frac{1}{\sqrt{(2\pi)^2 \sigma_i \sigma_j}}$$




$$O[i,j] = \sum_{\substack{k \\ k \leq K}} \sum_{\substack{l \\ l \leq L}} \underbrace{I[i+k, j+l]}_a \underbrace{w[k,l]}_b$$

$$\underline{a} = \begin{bmatrix} I[i,j] \\ I[i, j+1] \\ \vdots \\ I[i+K, j+L] \end{bmatrix} \in \mathbb{R}^{KL} \quad \underline{b} = \begin{bmatrix} w[0,0] \\ w[0,1] \\ \vdots \\ w[K,L] \end{bmatrix} \in \mathbb{R}^{KL}$$

$$O[i,j] = \underline{a}^T \underline{b} = |\underline{a}| |\underline{b}| \cos \theta$$



$$\underline{a} = \lambda \underline{b} \Rightarrow \underline{a} \propto \underline{b}$$



$|O[i,j]|$ is maximized

when $I[i:i+k, j:j+l] \propto w[0:k, 0:l]$

Template matching

$O[i,j]$ = similarity measure

maybe = - Distance measure


maybe = $\frac{1}{\text{Distance measure}}$

$\underline{a}^T \underline{b}$ is also called cosine similarity between \underline{a} and \underline{b}

cosine distance $(1 - \underline{a}^T \underline{b})$

$$O[i,j] = - \left(\sum_{k=0}^{K-1} \sum_{l=0}^{L-1} \left(I[i+k, j+l] - w[k,l] \right)^2 \right)$$

$$= - \|\underline{a} - \underline{b}\|^2 = \text{Sum of squares}$$



$$O_w = 32 - 3 + 1 + 2 \cdot 1 = 32$$

```
In [1]: # Adapted from: Chapter 7 and 8 of Deep Learning with Pytorch by Eli Stevens
try:
    import torch as t
    import torch.nn as tnn
except ImportError:
    print("Colab users: pytorch comes preinstalled. Select Change Ru")
    print("Local users: Please install pytorch for your hardware using instr")
    print("ACG users: Please follow instructions here: https://vikasdhiman.i")

    raise

if t.cuda.is_available():
    DEVICE="cuda"
elif t.mps.is_available():
    DEVICE="mps"
else:
    DEVICE="cpu"

DTYPE = t.get_default_dtype()
```

```
In [2]: ## Doing it the Pytorch way without using our custom feature extraction
```

```
import torch
import torch.nn
import torch.optim
import torchvision
from torchvision.transforms import ToTensor, Compose, Normalize
from torch.utils.data import DataLoader

torch.manual_seed(17)
DATASET_MEAN = [0.4914, 0.4822, 0.4465]
DATASET_STD = [0.2470, 0.2435, 0.2616]
# Getting the dataset, the Pytorch way
all_training_data = torchvision.datasets.CIFAR10(
    root="data",
    train=True,
    download=True,
    transform=Compose([ToTensor(),
                        Normalize(DATASET_MEAN, # dataset mean
                                DATASET_STD)]) # dataset std
)

test_data = torchvision.datasets.CIFAR10(
    root="data",
    train=False,
    download=True,
    transform=Compose([ToTensor(),
                        Normalize(DATASET_MEAN, # dataset mean
                                DATASET_STD)]) # dataset std
)
```

Files already downloaded and verified
Files already downloaded and verified

```
In [3]: training_data, validation_data = torch.utils.data.random_split(all_training_
```


```
In [4]: img, label = all_training_data[99]  
img.shape, label
```

```
Out[4]: (torch.Size([3, 32, 32]), 1)
```

```
In [5]: import matplotlib.pyplot as plt  
plt.imshow(img.permute(1, 2, 0))
```


Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```
Out[5]: <matplotlib.image.AxesImage at 0x7efdecbb5b0>
```



```
In [6]: plt.imshow((img.permute(1, 2, 0) * torch.Tensor(DATASET_STD)  
+ torch.Tensor(DATASET_MEAN)))
```

```
Out[6]: <matplotlib.image.AxesImage at 0x7efdecbb04940>
```



```
In [7]: imgs = torch.stack([img_t for img_t, _ in all_training_data], dim=3)
        imgs.reshape(3, -1).mean(dim=-1), imgs.reshape(3, -1).std(dim=-1)
```

```
Out[7]: (tensor([-1.2762e-06, -1.7074e-04,  1.1819e-04]),
         tensor([1.0001, 0.9999, 1.0000]))
```

```
In [8]: import pickle
        cifar_meta = pickle.load(open("data/cifar-10-batches-py/batches.meta", "rb"))
        class_names = [c.decode('utf-8') for c in cifar_meta[b'label_names']]
        class_names
```

```
Out[8]: ['airplane',
         'automobile',
         'bird',
         'cat',
         'deer',
         'dog',
         'frog',
         'horse',
         'ship',
         'truck']
```

```
In [9]: # Hyper parameters
        learning_rate = 1e-3 # controls how fast the gradient descent goes
        batch_size = 64
        epochs = 5
        momentum = 0.9

        training_dataloader = DataLoader(training_data, shuffle=True, batch_size=batch_size)
        validation_dataloader = DataLoader(validation_data, batch_size=batch_size)
        test_dataloader = DataLoader(test_data, batch_size=batch_size)
```

```
X, y = next(iter(training_dataloader))  
X.shape
```

Out[9]: torch.Size([64, 3, 32, 32])

```
In [10]: !pip install tensorboard
```


Requirement already satisfied: tensorboard in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (2.12.0)

Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from tensorboard) (0.4.6)

Requirement already satisfied: grpcio>=1.48.2 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from tensorboard) (1.53.0)

Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from tensorboard) (0.7.0)

Requirement already satisfied: requests<3,>=2.21.0 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from tensorboard) (2.28.1)

Requirement already satisfied: markdown>=2.6.8 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from tensorboard) (3.4.3)

Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from tensorboard) (1.8.1)

Requirement already satisfied: google-auth<3,>=1.6.3 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from tensorboard) (2.17.0)

Requirement already satisfied: wheel>=0.26 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from tensorboard) (0.37.1)

Requirement already satisfied: setuptools>=41.0.0 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from tensorboard) (65.6.3)

Requirement already satisfied: werkzeug>=1.0.1 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from tensorboard) (2.2.3)

Requirement already satisfied: numpy>=1.12.0 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from tensorboard) (1.23.5)

Requirement already satisfied: protobuf>=3.19.6 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from tensorboard) (4.22.1)

Requirement already satisfied: absl-py>=0.4 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from tensorboard) (1.4.0)

Requirement already satisfied: rsa<5,>=3.1.4 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from google-auth<3,>=1.6.3->tensorboard) (4.9)

Requirement already satisfied: cachetools<6.0,>=2.0.0 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from google-auth<3,>=1.6.3->tensorboard) (5.3.0)

Requirement already satisfied: pyasn1-modules>=0.2.1 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from google-auth<3,>=1.6.3->tensorboard) (0.2.8)

Requirement already satisfied: six>=1.9.0 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from google-auth<3,>=1.6.3->tensorboard) (1.16.0)

Requirement already satisfied: requests-oauthlib>=0.7.0 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from google-auth-oauthlib<0.5,>=0.4.1->tensorboard) (1.3.1)

Requirement already satisfied: certifi>=2017.4.17 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from requests<3,>=2.21.0->tensorboard) (2022.12.7)

Requirement already satisfied: urllib3<1.27,>=1.21.1 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from requests<3,>=2.21.0->tensorboard) (1.26.14)

Requirement already satisfied: charset-normalizer<3,>=2 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from requests<3,>=2.21.0->tensorboard) (2.0.4)
Requirement already satisfied: idna<4,>=2.5 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from requests<3,>=2.21.0->tensorboard) (3.4)
Requirement already satisfied: MarkupSafe>=2.1.1 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from werkzeug>=1.0.1->tensorboard) (2.1.1)
Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from pyasn1-modules>=0.2.1->google-auth<3,>=1.6.3->tensorboard) (0.4.8)
Requirement already satisfied: oauthlib>=3.0.0 in /home/vdhiran/.local/miniconda3/envs/ece490/lib/python3.10/site-packages (from requests-oauthlib>=0.7.0->google-auth-oauthlib<0.5,>=0.4.1->tensorboard) (3.2.2)

In [17]: `%tensorboard --logdir=runs`

Reusing TensorBoard on port 6006 (pid 543069), started 0:05:49 ago. (Use '!kill 543069' to kill it.)

```
In [22]: from torch.utils.tensorboard import SummaryWriter
import os
writer = SummaryWriter()

loss = torch.nn.CrossEntropyLoss()
# TODO:
# Define model = ?
```

```

model = tnn.Sequential(
    torch.nn.Flatten(),
    tnn.Linear(3*32*32, 100),
    tnn.ReLU(),
    tnn.Linear(100, 10))

# Define optimizer
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate, momentum=momentum)

def loss_and_accuracy(model, loss, validation_dataloader, device=DEVICE):
    # Validation loop
    validation_size = len(validation_dataloader.dataset)
    num_batches = len(validation_dataloader)
    test_loss, correct = 0, 0

    with torch.no_grad():
        model.eval() # Put model in eval mode, affects layers like dropout
        for X, y in validation_dataloader:
            X = X.to(device)
            y = y.to(device)
            pred = model(X)
            test_loss += loss(pred, y)
            correct += (pred.argmax(dim=-1) == y).type(DTYPE).sum()

    test_loss /= num_batches
    correct /= validation_size
    return test_loss, correct

def train(model, loss, training_dataloader, validation_dataloader, device=DEVICE):
    model.to(device)
    t0 = 0
    if os.path.exists("runs/model_ckpt.pt"):
        checkpoint = torch.load("runs/model_ckpt.pt")
        model.load_state_dict(checkpoint['model_state_dict'])
        optimizer.load_state_dict(checkpoint['optimizer_state_dict'])
        t0 = checkpoint['epoch']

    for t in range(t0, epochs):
        # Train loop
        training_size = len(training_dataloader.dataset)
        nbatches = len(training_dataloader)
        model.train() # Put model in train mode, affects layers like dropout
        for batch, (X, y) in enumerate(training_dataloader):
            X = X.to(device)
            y = y.to(device)
            # Compute prediction and loss
            pred = model(X)
            loss_t = loss(pred, y)

            # Backpropagation
            optimizer.zero_grad()
            loss_t.backward()
            optimizer.step()

            if batch % 100 == 0:

```

```

        writer.add_scalar("Train/loss_batch", loss_t, t*nbatches +
        loss_t, current = loss_t.item(), (batch + 1) * len(X)
        print(f"loss: {loss_t:>7f}  [{current:>5d}/{training_size:>5d}]")

    writer.add_scalar("Train/loss", loss_t, t)
    valid_loss, correct = loss_and_accuracy(model, loss, validation_data_loader)
    writer.add_scalar("Valid/loss", valid_loss, t)
    writer.add_scalar("Valid/accuracy", correct, t)
    print(f"Validation Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss: {valid_loss:>0.1f}")
    if t % 3 == 0:
        torch.save({
            'epoch': t,
            'model_state_dict': model.state_dict(),
            'optimizer_state_dict': optimizer.state_dict()
        }, "runs/model_ckpt.pt")
    return model

trained_model = train(model, loss, training_data_loader, validation_data_loader)

test_loss, correct = loss_and_accuracy(model, loss, test_data_loader)
print(f"Test Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss: {test_loss:>0.1f}")

Validation Error: 4864/45000]
Accuracy: 48.4%, Avg loss: 1.453683

Validation Error: 4864/45000]
Accuracy: 49.2%, Avg loss: 1.440753

Test Error:
Accuracy: 50.3%, Avg loss: 1.427571

```

Limits of fully connected layers

Convolutions bring translation invariance

```

In [32]: conv = torch.nn.Conv2d(3, 1, kernel_size=3)
conv

```

```

Out[32]: Conv2d(3, 1, kernel_size=(3, 3), stride=(1, 1))

```

```


In [33]: img, label = training_data[99]
plt.imshow(img.permute(1, 2, 0) * t.Tensor(DATASET_STD) + t.Tensor(DATASET_MEAN))

```

```

Out[33]: <matplotlib.image.AxesImage at 0x7efd51214ee0>

```



Learning with Convolutions

Maxpooling

In []:

